

## **CSE 508 Course Project Presentation**

# **nftables on Android**

Team Members

Subhasis Das

Gaurav Piyush

Chaitanya Saxena

# Tasks Attempted

- 60 Points
  - Port nftables
  - Port the user space libraries in android
  - Port required kernel modules (apply patches)
  - Table and Input chains
  - Add / delete rules
  - Drop packets/Block traffic
- 20 Points : Redirection and Source Spoofing
  - Using NAT module
  - Create NAT table and new NAT chain
  - Add NAT rules
  - Testing
- Packet Selector
  - Matching Transport protocol , IPV4/IPV6 Headers
  - Matching TCP/UDP/UDPlite traffic
  - Matching Sender/Receiver
  - Matching packet meta information

# nftables

- Common platform for iptables, etables, ip6table, arptable (aims to replace them)
- Protocol dependency in user space.
- Core common hooks in kernel space.
- Common language for rule generation and parsing.

# Goldfish

- Android OS : Built upon linux kernel 3.4
- Does not support nftables
- The Android emulator runs a virtual CPU that Google calls Goldfish.  
Goldfish executes ARM926T instructions and has hooks for input and output ([2])
- Kernel Version 3.10

# User Space Libraries

Following user level modules were successfully ported to the kernel :

libmnl

libnftnl

libgmp

libreadline

libncurses

nftable

- Set up the proper environment path
- Check cross compilation
- Test each library by using a user program



gaurav@gaurav-Studio-1555: ~... x gau

bugreports  
dalvik-cache  
data  
dontpanic  
drm  
libgmp\_test  
libmnl\_test  
libnftnl\_test  
libreadline\_test  
local  
lost+found  
media  
mediadm  
misc  
nativebenchmark  
nativetest  
property  
resource-cache  
security  
ssh  
system  
user

Options:

-h/--help

Show this help

-v/--version

Show version information

-f/--file <filename>

Read input from <filename>

-i/--interactive

Read input from interactive CLI

-n/--numeric

When specified once, show network addresses numerically.

When specified twice, also show Internet services,  
user IDs and group IDs numerically.

When specified thrice, also show protocols numerically.

-a/--handle

Output rule handle.

-I/--includepath <directory>

Add <directory> to the paths searched for include files.

--debug <level [,level...]>

Specify debugging level (scanner, parser, eval, netlink, mnl, proto-ctx, segtree, all)

root@generic\_x86:/data # ./libmnl\_test  
TEST  
4096  
DONE

root@generic\_x86:/data # ./libnftnl\_test  
libnftnl7

root@generic\_x86:/data # ./libgmp\_test

7612058254738945

\*

9263591128439081

70514995317761165008628990709545

root@generic\_x86:/data # ./libreadline\_test

READLINE6

DONERoot@generic\_x86:/data #

# Kernel Space Patches

1. Patch 1 - 96518518cc417bb0a8c80b9fb736202e28acdf96
  - a. Core implementation for nftables in kernel space
  - b. Storage of rule list per chain - new private data pointer
2. Patch 2 - f59cb0453cd885736daa11ae2445982c5ab2fc83
  - a. Creation of common module remove duplication of code for iptable and nftable (nat\_decode\_session, alloc\_null\_binding)
3. Patch 3 - 795aa6ef6a1aba99050735eadd0c2341b789b53b
  - a. The user space nftable utility communicates to the kernel space through hooks. This patch creates a generic hook function consisting of common hook functions
4. Patch 4 - 20a69341f2d00cd042e81c82289fba8a13c05a25
  - a. Defines nftable sets, different from rule sets.
  - b. Defines operations like creation, deletion, lookup etc. on sets.
  - c. Defines lockless operation on sets if defined as a constant (not allowed to change when a rule is linked).

# NFT rule to drop packet

```
PING 127.0.0.1 (127.0.0.1) 56(84) bytes of data.  
64 bytes from 127.0.0.1: icmp_seq=1 ttl=64 time=0.063 ms  
64 bytes from 127.0.0.1: icmp_seq=2 ttl=64 time=0.075 ms  
64 bytes from 127.0.0.1: icmp_seq=3 ttl=64 time=0.082 ms  
^C  
--- 127.0.0.1 ping statistics ---  
3 packets transmitted, 3 received, 0% packet loss, time 1998ms  
rtt min/avg/max/mdev = 0.063/0.073/0.082/0.010 ms
```

```
nft add rule ip filter output ip daddr 127.0.0.1 drop
```

```
PING 127.0.0.1 (127.0.0.1) 56(84) bytes of data.  
ping: sendmsg: Operation not permitted  
ping: sendmsg: Operation not permitted  
ping: sendmsg: Operation not permitted  
ping: sendmsg: Operation not permitted  
ping: sendmsg: Operation not permitted  
^C  
--- 127.0.0.1 ping statistics ---  
5 packets transmitted, 0 received, 100% packet loss, time 4016ms
```

Add Rule to drop  
Packets



Packets Dropped



# Rule Chain

```
nft list table filter
```

```
table ip filter {  
    chain input {  
        type filter hook input priority 0;  
    }  
  
    chain forward {  
        type filter hook forward priority 0;  
    }  
  
    chain output {  
        type filter hook output priority 0;  
        ip daddr 1.2.3.4 drop  
        ip daddr localhost drop  
    }  
}
```

```
nft list table filter -a
```

```
table ip filter {  
    chain input {  
        type filter hook input priority 0;  
    }  
  
    chain forward {  
        type filter hook forward priority 0;  
    }  
  
    chain output {  
        type filter hook output priority 0;  
        ip daddr 1.2.3.4 drop # handle 4  
        ip daddr localhost drop # handle 5  
    }  
}
```

## Next : Delete Rule

```
nft delete rule filter output handle 5
```

# NFT more rules to drop packets

- Add rule to a network :  
`nft add rule ip filter output ip daddr 172.24.241.0/24 counter`
- Add rule to a port 80  
`nft add rule ip filter input tcp dport 80 drop`
- A combined rule ( filters ICMP and drops O/P to destination)  
`nft add rule ip filter output ip protocol icmp ip daddr 127.0.0.1 counter drop`

# Redirection

- Port nat module
- Step 1: Make kernel aware of NAT

*modprobe nft\_chain\_nat\_ipv4*

- Step 2: Create NAT dedicated chains

- `sudo nft add table nat`
- `sudo nft add chain nat post '{ type nat hook postrouting priority 0 \; \; }`
- `sudo nft add chain nat pre '{ type nat hook prerouting priority 0 \; \; }`

- Step 3 : Add some nat rules

`nft add rule nat pre udp dport 53 ip saddr 192.168.56.0/24 dnat 8.8.8.8:53`

(Redirects all DNS traffic from 192.168.56.0/24 to the 8.8.8.8(Google Public DNS))

```
table ip nat {  
    chain post {  
        type nat hook postrouting priority 0;  
    }  
  
    chain pre {  
        type nat hook prerouting priority 0;  
        udp dport domain ip saddr 192.168.56.0/24 dnat google-public-dn  
s-a.google.com:domain  
    }  
}
```

# Packet Selectors and actions

- Matching Transport protocol

```
nft add rule filter output ip protocol tcp
```

- Matching IPV4 heading : Sender and Receiver

```
nft add rule filter input ip saddr 192.168.1.100 ip daddr 192.168.1.1 counter
```

- Matching TCP traffic : matches and drops all tcp traffic for low TCP ports (1-1024)

```
nft add rule filter input tcp dport 1-1024 counter drop
```

- Matching traffic based on user name

```
nft add rule filter output meta skuid 1000 counter
```

# Packet Selectors and Action

```
table ip filter {  
    chain input {  
        type filter hook input priority 0;  
    }  
  
    chain forward {  
        type filter hook forward priority 0;  
    }  
  
    chain output {  
        type filter hook output priority 0;  
        ip daddr localhost drop  
        skuid ron counter packets 8 bytes 528  
        skuid ron counter packets 8 bytes 528  
    }  
}
```

# Challenges Faced

- No proper documentation available about nftable porting
- Cross compilation issues
- Locating and adding kernel patches
- Running the emulator
- Running internet on emulator

# References

1. <http://en.wikipedia.org/wiki/Nftables> : Netfilter Introduction
2. <https://groups.google.com/forum/#!topic/android-kernel/M4SjXuUeUo> : Goldfish
3. <https://wiki.archlinux.org/index.php/Nftables> : A good basic documentation on usage and design of nftables.
4. <https://github.com/sam8dec/NetSec> : We have referred to this excellent write up by Samudra for our initial setup.
5. [http://en.wikipedia.org/wiki/GNU\\_Multiple\\_Precision\\_Arithmetic\\_Library#Example](http://en.wikipedia.org/wiki/GNU_Multiple_Precision_Arithmetic_Library#Example) : Excellent reference for writing test cases
6. Source for nfnetlink\_compat.h : [https://git.netfilter.org/libnetfilter\\_acct/tree/include/linux/netfilter](https://git.netfilter.org/libnetfilter_acct/tree/include/linux/netfilter).
7. [http://kernelnewbies.org/nftables\\_examples](http://kernelnewbies.org/nftables_examples) : Excellent examples of NAT rule handling
8. Links to patches: [[Patch 1](#)] [[Patch 2](#)] [[Patch 3](#)] [[Patch 4](#)]